

# *Improving Query Performance*

---

陆伟

College of Software

**Database Systems**

---

May 19, 2021

# Outline

---

- Introduction
- Database Storage and Query Performance
- Hash Files
- Indexes
- Query Optimization and Index Choice

# Introduction

---

- A DBMS stores a database on a secondary storage, which typically is a disk. As we have learned in our discussion of database recovery, disks are extremely slow devices compared to main memory.
- So ...


# Introduction

---

- Attempting to improve the performance in a database system
  - Minimize the data flow between the main memory and the disk. (buffer)
  - Minimize the number of disk accesses required for locating a row. (access paths)
    - File organization
    - Access algorithms
- **Indexes** are examples of access paths and can greatly reduce the time it takes to process a query.

# Database Storage and Query Performance

- Hard Disk Design
- Measuring Hard Disk Performance
  - Seek time: 8 to 14 milliseconds
  - Rotational latency
  - Block transfer time



# Database Storage and Query Performance

---

- File or physical organization and Access Methods
  - The arrangements of the data into records (or tuples, rows, data items, objects) and pages (or blocks) on a disk is called **file or physical organization**.
  - Each file organization supports specific methods to store and retrieve records from the file.
  - Examples:
    - ordered files - binary search algorithm
    - heap files - linear search

# Database Storage and Query Performance

---

- Improving Query Performance with Indexes
  - About blocking factor
  - Example

Consider a table consisting of 100,000 records and a blocking factor of 100. Sequentially scanning this entire table for a record can require as many as 1,000 disk blocks being read.

Assuming you have an index which is stored on a single disk block, on the other hand, may require as few as two disk blocks being read to find a record in this table (one disk block read to load the index and one disk block read to load the data).

# Hash Files

---

- Hashing: Basic Organization
- About Collisions
- Properties of a Good Hash Function
  - It should be computed easily and efficiently.
  - It should minimize the number of collisions by spreading keys around the file as evenly and uniformly as possible.
- Pros and Cons of Hashing



# Indexes

---

- An index is an **auxiliary file** that makes it more efficient to search for a record in a data file. It is usually defined on one field of the data file, called the **indexing field**.
- Each **index entry** is a pairing of an indexing value with a pointer to the page in which the value appears as part of a record.
- For fast search, the **entries are sorted by the indexing field values**.

# Indexes

---

- According to the different properties of the indexing fields involved, indexes can be classified as
  - Primary Index
  - Secondary Index
  - Cluster Index
- Indexes are also classified as
  - Sparse
  - Dense

# Indexes

---

- Indexed Sequential Access Method (ISAM)
  - An ISAM is a sparse, primary index.
  - It includes one index entry for each page in the data file.
  - The first record in a data page is called the page anchor.
  - The field value of an index entry is the key value of the page anchor.

# Indexes

---

- Cluster Index
- Secondary ISAM
- Multi-Level ISAM


# Indexes

---

- Drawbacks of ISAM
  - ISAMs are static structures and can offer great performance guarantees for searching, as long as insertions, deletions, and updates are very infrequent.
  - But in dynamic database system environments, ...

# Indexes

- B+ Trees



# Indexes

---

- Indexes in SQL
  - Most existing DBMSs use fixed physical design and automatically build the best possible access path, which usually is some B+ tree.
- The general syntax of the CREATE INDEX statement

```
CREATE [UNIQUE] INDEX index-name  
    ON table (indexing-attribute-order-list) [CLUSTER];
```

# Query Optimization and Index Choice

---

- Query optimization is a process by which **programmers** and **database administrators** make use of access methods to ensure that an application accesses data as efficiently as possible.
  - One of the steps in optimizing queries is to **examine the physical storage** characteristics of tables accessed by those queries.
  - The next step in query optimization is analyzing the run-time characteristics of queries. (**Find bottlenecks** )



# Query Optimization and Index Choice

- PostgreSQL Analysis Tools
  - The **pg\_class** table stores information and statistics about each table

```
SELECT relname, reltuples, relpages FROM  
pg_class WHERE relname = 'student';
```

relname		reltuples		relpages
-----	+	-----	+	-----
student		249999		4425

# Query Optimization and Index Choice

---

- PostgreSQL Analysis Tools
  - It is important for the query planner to have accurate statistics in the **pg\_class** table.
  - The analyze command updates the statistics in the **pg\_class** table.
  - PostgreSQL does not update the data stored in the **pg\_class** table for each update to a user table.
  - To allow the query planner to make the best possible plan, the database administrator should run the analyze command often.

# Query Optimization and Index Choice

- PostgreSQL Analysis Tools
  - PostgreSQL can report a query's execution plan using **explain** command.

```
explain select distinct course_id from course where course_term = 'Fal02';
```

NOTICE: QUERY PLAN:

Unique (cost=12223.09..12339.76 rows=4667 width=4)

-> Sort (cost=12223.09..12223.09 rows=46666 width=4)

-> Seq Scan on course (cost=0.00..8279.99 rows=46666 width=4)

# Query Optimization and Index Choice

---

- PostgreSQL Analysis Tools
  - It is important to note that the explain command does not actually execute a query, it only gets the plan for a query's execution and the **estimated cost** of running the query from PostgreSQL's query planner.

# Query Optimization and Index Choice

---

- PostgreSQL Analysis Tools
  - The **explain analyze** command will execute a query and display the plan of execution for a query as well as the actual run-time in milliseconds and the actual number of disk page fetches for a query.

# Query Optimization and Index Choice

## ■ PostgreSQL Analysis Tools

```
explain analyze select distinct course_id  
from course  
where course_term = 'Fal02';
```

NOTICE: QUERY PLAN:

```
Unique (cost=12223.09..12339.76 rows=4667 width=4)  
  (actual time=1643.87..1797.34 rows=41803 loops=1)  
    -> Sort (cost=12223.09..12223.09 rows=46666 width=4)  
      (actual time=1643.86..1706.05 rows=41803 loops=1)  
        -> Seq Scan on course (cost=0.00..8279.99 rows=46666 width=4)  
          (actual time=184.83..1075.11 rows=41803 loops=1)  
Total runtime: 1899.24 msec
```

# Query Optimization and Index Choice

---

- PostgreSQL Indexes
  - PostgreSQL supports both **hash** and **b-tree** indexes. PostgreSQL also supports **clustering** indexes.

```
create index index_name on table using  
index_type(column[,column[,...]]);
```

```
cluster index_name on table;
```

# Query Optimization and Index Choice

---

- PostgreSQL Indexes

```
create index assessment_point_index on assessment  
using btree(total_points);
```

```
cluster assessment_point_index on assessment;
```

```
create index assessment_id_index on assessment  
using hash(assessment_id);
```



# Query Optimization and Index Choice

---

- Application-Readings
  - Exercise

# Query Optimization and Index Choice

---

- Index Choice and Optimization Heuristics

# Summary

---



- In this chapter you should have learned:
  - Factors affecting query performance
  - Principles and types of index
  - Query Optimization and Index Choice

# Questions

---

